

Please see the **API and JavaDoc** for the most current version of the spec.

JSR107: JCACHE Java™ Caching API

JSR107 Expert Group

Specifications Leads:
Greg Luck, Terracotta Inc.
Yannis Cosmadopoulos, Oracle Corporation

Early Draft 1

Comments to: jsr107@googlegroups.com

Specification: JCACHE - Java™ Temporary Caching API("Specification")
Version: Early Draft 1
Status: In progress
Release: Unreleased

Copyright 2011 Oracle
Copyright 2011 Greg Luck
All rights reserved.

NOTICE

The Specification is protected by copyright and the information described therein may be protected by one or more U.S. patents, foreign patents, or pending applications. Except as provided under the following license, no part of the Specification may be reproduced in any form by any means without the prior written authorization of Sun Microsystems, Inc. ("Sun") and its licensors, if any. Any use of the Specification and the information described therein will be governed by the terms and conditions of this Agreement.

Subject to the terms and conditions of this license, including your compliance with Paragraphs 1 and 2 below, Sun hereby grants you a fully-paid, non-exclusive, non-transferable, limited license (without the right to sublicense) under Sun's intellectual property rights to:

1. Review the Specification for the purposes of evaluation. This includes: (i) developing implementations of the Specification for your internal, non-commercial use; (ii) discussing the Specification with any third party; and (iii) excerpting brief portions of the Specification in oral or written communications which discuss the Specification provided that such excerpts do not in the aggregate constitute a significant portion of the Technology.
2. Distribute implementations of the Specification to third parties for their testing and evaluation use, provided that any such implementation:
 - (i) does not modify, subset, superset or otherwise extend the Licensor Name Space, or include any public or protected packages, classes, Java interfaces, fields or methods within the Licensor Name Space other than those required/authorized by the Specification or Specifications being implemented; (ii) is clearly and prominently marked with the word "UNTESTED" or "EARLY ACCESS" or "INCOMPATIBLE" or "UNSTABLE" or "BETA" in any list of available builds and in proximity to every link initiating its download, where the list or link is under Licensees control; and
 - (iii) includes the following notice: "This is an implementation of an early-draft specification developed under the Java Community Process (JCP) and is made available for testing and evaluation purposes only. The code is not compatible with any specification of the JCP."

The grant set forth above concerning your distribution of implementations of the specification is contingent upon your agreement to terminate development and distribution of your early draft implementation as soon as feasible following final completion of the specification. If you fail to do so, the foregoing grant shall be considered null and void.

No provision of this Agreement shall be understood to restrict your ability to make and distribute to third parties applications written to the Specification.

Other than this limited license, you acquire no right, title or interest in or to the Specification or any other Sun intellectual property, and the Specification may only be used in accordance with the license terms set forth herein. This license will expire on the earlier of: (a) two (2) years from the date of Release listed above; (b) the date on which the final version of the Specification is publicly released; or (c) the date on which the Java Specification Request (JSR) to which the Specification corresponds is withdrawn. In addition, this license will terminate immediately without notice from Sun if you fail to comply with any provision of this license. Upon termination, you must cease use of or destroy the Specification.

"Licensor Name Space" means the public class or interface declarations whose names begin with "java", "javax", "com.sun" or their equivalents in any subsequent naming convention adopted by Sun through the Java Community Process, or any recognized successors or replacements thereof

TRADEMARKS

No right, title, or interest in or to any trademarks, service marks, or trade names of Sun or Sun's licensors is granted hereunder. Sun, Sun Microsystems, the Sun logo, Java, are trademarks or registered trademarks of Sun Microsystems, Inc. in the U.S. and other countries.

DISCLAIMER OF WARRANTIES

THE SPECIFICATION IS PROVIDED "AS IS" AND IS EXPERIMENTAL AND MAY CONTAIN DEFECTS OR DEFICIENCIES WHICH CANNOT OR WILL NOT BE CORRECTED BY SUN. SUN MAKES NO

REPRESENTATIONS OR WARRANTIES, EITHER EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO, WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, OR NON-INFRINGEMENT THAT THE CONTENTS OF THE SPECIFICATION ARE SUITABLE FOR ANY PURPOSE OR THAT ANY PRACTICE OR IMPLEMENTATION OF SUCH CONTENTS WILL NOT INFRINGE ANY THIRD PARTY PATENTS, COPYRIGHTS, TRADE SECRETS OR OTHER RIGHTS. This document does not represent any commitment to release or implement any portion of the Specification in any product.

THE SPECIFICATION COULD INCLUDE TECHNICAL INACCURACIES OR TYPOGRAPHICAL ERRORS. CHANGES ARE PERIODICALLY ADDED TO THE INFORMATION THEREIN; THESE CHANGES WILL BE INCORPORATED INTO NEW VERSIONS OF THE SPECIFICATION, IF ANY. SUN MAY MAKE IMPROVEMENTS AND/OR CHANGES TO THE PRODUCT(S) AND/OR THE PROGRAM(S) DESCRIBED IN THE SPECIFICATION AT ANY TIME. Any use of such changes in the Specification will be governed by the then-current license for the applicable version of the Specification.

LIMITATION OF LIABILITY

TO THE EXTENT NOT PROHIBITED BY LAW, IN NO EVENT WILL SUN OR ITS LICENSORS BE LIABLE FOR ANY DAMAGES, INCLUDING WITHOUT LIMITATION, LOST REVENUE, PROFITS OR DATA, OR FOR SPECIAL, INDIRECT, CONSEQUENTIAL, INCIDENTAL OR PUNITIVE DAMAGES, HOWEVER CAUSED AND REGARDLESS OF THE THEORY OF LIABILITY, ARISING OUT OF OR RELATED TO ANY FURNISHING, PRACTICING, MODIFYING OR ANY USE OF THE SPECIFICATION, EVEN IF SUN AND/OR ITS LICENSORS HAVE BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

You will hold Sun (and its licensors) harmless from any claims based on your use of the Specification for any purposes other than the limited right of evaluation as described above, and from any claims that later versions or releases of any Specification furnished to you are incompatible with the Specification provided to you under this license.

RESTRICTED RIGHTS LEGEND

If this Software is being acquired by or on behalf of the U.S. Government or by a U.S. Government prime contractor or subcontractor (at any tier), then the Government's rights in the Software and accompanying documentation shall be only as set forth in this license; this is in accordance with 48 C.F.R. 227.7201 through 227.7202-4 (for Department of Defense (DoD) acquisitions) and with 48 C.F.R. 2.101 and 12.212 (for non-DoD acquisitions).

REPORT

You may wish to report any ambiguities, inconsistencies or inaccuracies you may find in connection with your evaluation of the Specification ("Feedback"). To the extent that you provide Sun with any Feedback, you hereby: (i) agree that such Feedback is provided on a non-proprietary and non-confidential basis, and (ii) grant Sun a perpetual, non-exclusive, worldwide, fully paid-up, irrevocable license, with the right to sublicense through multiple levels of sublicensees, to incorporate, disclose, and use without limitation the Feedback for any purpose related to the Specification and future versions, implementations, and test suites thereof.

GENERAL TERMS

Any action related to this Agreement will be governed by California law and controlling U.S. federal law. The U.N. Convention for the International Sale of Goods and the choice of law rules of any jurisdiction will not apply.

The Specification is subject to U.S. export control laws and may be subject to export or import regulations in other countries. Licensee agrees to comply strictly with all such laws and regulations and acknowledges that it has the responsibility to obtain such licenses to export, re-export or import as may be required after delivery to Licensee.

This Agreement is the parties' entire agreement relating to its subject matter. It supersedes all prior or contemporaneous oral or written communications, proposals, conditions, representations and warranties and prevails over any conflicting or additional terms of any quote, order, acknowledgment, or other communication between the parties relating to its subject matter during the term of this Agreement. No modification to this Agreement will be binding, unless in writing and signed by an authorized representative of each party.

Contents

[Chapter 1 - Introduction](#)

- [1.1 Status](#)
- [1.2 Purpose](#)
- [1.3 Goals](#)
- [1.4 Conventions](#)
- [1.5 Expert Group Members](#)
- [1.6 Acknowledgements](#)

[Chapter 2 - Caches](#)

- [2.1 Requirements of Keys and Values](#)
- [2.3 Generics](#)
- [2.4 Looking up a Cache](#)
- [2.5 Package](#)
- [2.6 Lifecycle](#)
- [2.7 Dependency on CacheManager](#)

[Chapter 3 Cache Configuration](#)

- [3.2 Resource Limits](#)

[Chapter 3 - Cache Operations](#)

- [Cache Operations](#)
- [Cache Entry Expiration](#)
- [Read-Through Caching](#)
- [Write-Through Caching](#)
- [CacheLoader](#)
- [CacheEntryListener](#)
- [Classloaders](#)

[Chapter 4 - Cache Managers](#)

- [Singleton versus Instance](#)
- [Singleton Construction](#)
- [Instance Construction](#)
- [Default CacheManager](#)
- [Caches must be in a CacheManager](#)

[Chapter 5 - Transactions](#)

- [All or nothing](#)
- [XA Transactions](#)
- [Enlistment](#)
- [Recovery](#)
- [Local Transactions](#)
- [Recovery](#)

[Chapter 6 - Isolation Levels](#)

- [READ_COMMITTED](#)
- [READ_UNCOMMITTED](#)
- [SERIALIZABLE](#)
- [REPEATABLE_READ](#)

[Chapter 7 - Caching Annotations](#)

- [7.1 - Annotations Overview](#)
- [7.1.1 - Annotation Inheritance and Ordering](#)
- [7.1.3 - Transactions](#)
- [7.2 - Annotations](#)
- [7.2.1 - @CacheDefaults](#)
- [Code Example 7-1](#) [@CacheDefaults Annotation Example](#)
- [7.2.2 - @CacheResult](#)
- [7.2.3 - @CachePut](#)
- [7.2.4 - @CacheRemoveEntry](#)
- [7.2.5 - @CacheRemoveAll](#)

- [7.2.5 - @CacheKeyParam](#)
- [7.2.6 - @CacheValue](#)
- [7.3 - Cache Resolution](#)
 - [7.3.1 - Cache Name](#)
 - [7.3.2 - CacheResolverFactory](#)
 - [7.3.3 - CacheResolver](#)
- [7.4 - Key Generation](#)
- [7.5 - Annotation Support Classes](#)
 - [7.5.1 - CacheMethodDetails](#)
 - [7.5.2 - CacheInvocationContext](#)
 - [7.5.3 - CacheKeyInvocationContext](#)
 - [7.5.4 - CacheInvocationParameter](#)
 - [7.5.5 - CacheKey](#)
- [Chapter 8 - Container and Provider Contracts for Deployment and Bootstrapping](#)
 - [8.2 Use of Caching](#)
- [Chapter 9 - Glossary](#)
- [Chapter 10 - Bibliography](#)
- [Appendix A - Revision History](#)
 - [Early Draft 1](#)
- [Appendix B - Open Issues](#)
 - [What is a Read-Through Cache?](#)
 - [Option 1 - Cache as Proxy](#)
 - [Option 2 - Cache is not a proxy - limited read-through operation](#)
 - [What is a write-through Cache?](#)
- [Statistics effects](#)
- [Listeners](#)
- [Project Management](#)
 - [Resources Needed](#)
 - [People](#)
 - [Tools](#)

Chapter 1 - Introduction

This document is the specification of the Java API for the temporary caching. The technical objective of this work is to provide a caching facility for the Java application developer.

This specification describes an API for caching, and an SPI so providers can implement their own caches.

Leading experts throughout the entire Java community have come together to build this Java caching standard.

1.1 Status

This is early draft 1, still in progress.

The latest API can be found online at:

<https://github.com/jsr107/jsr107spec>

The reference implementation can be obtained from:

<https://github.com/jsr107/RI>

Finally the TCK can be obtained from:

<https://github.com/jsr107/jsr107tck>

The expert group seeks feedback from the community on any aspect of this specification, please send comments to:

jsr-107-comments@jcp.org

or, for a publicly readable forum to:

jsr107@googlegroups.com

1.2 Purpose

Caching is a tried and true method for dramatically speeding up applications. Applications often use temporary data which is expensive to create, but has a lifetime over which it can be reused. For example, a servlet might create a web page from data obtained from multiple databases, network connections, and expensive computations; the data sets might be reusable over the same or different periods of time.

This specification standardizes caching of Java objects in a way that allows an efficient implementation, and removes from the programmer the burden of implementing cache expiration, mutual exclusion, spooling, and cache consistency.

1.3 Goals

The goals of the API are:

Object Cache The API will cache Java objects.

Support for By-Value caching and optionally, By-Reference. In the latter references are cached on heap

and in the former both keys and values are transformed into values.

Support for Flexible Implementations The specification will support both in-process and distributed implementations.

Java SE The specification will work with Java SE.

Java EE The specification will work with Java EE. This specification is targeted at inclusion in Java EE 7.

Annotations The specification will define optional runtime cache annotations.

Transactions Optional support for transactions, both local and XA will be defined.

1.4 Conventions

The regular Times font is used for information that is normative for this specification.

The italic Times font is used for paragraphs that contain non-normative information, such as notes describing typical use, or notes clarifying the text with prescriptive specification.

The Courier New font is used for code examples.

In addition, the keywords ‘MUST’, ‘MUST NOT’, ‘REQUIRED’, ‘SHALL’, ‘SHALL NOT’, ‘SHOULD’, ‘SHOULD NOT’, ‘RECOMMENDED’, ‘MAY’, and ‘OPTIONAL’ in this document are to be interpreted as described in RFC 2119.

Java code and sample data fragments are formatted as shown in figure 1.1:

```
package com.example.hello;

public class Hello {
    public static void main(String args[] {
        System.out.println("Hello Worlds");
    }
}
```

Figure 1: Example Java Code

1.5 Expert Group Members

This specification is being developed under the Java Community Process v2.7.

The following are active expert group members:

- Greg Luck
- Cameron Purdy, Oracle
- Yannis Cosmadopoulos, Oracle
- Manik Surtani, Red Hat, Inc.
- Nikita Ivanov, Grid Gain
- Chris Berry
- Andy Piper

- Jon Stevens
- Rick Hightower
- Ben Cotton
- David Mossakowski, Citigroup
- Bongjae Chang
- Steve Millidge
- Gabe Montero, IBM
- Eric Dalquist
- Pete Muir, Red Hat, Inc.
- William Newport, Goldman Sachs
- Ryan Gardner, Dealers.com
- Steve Harris, Terracotta

The following are members of JSR342 (Java EE 7) and are official observers of JSR107 which means inter alia they get read only access to the mailing list :

- Linda Demichiel, Oracle
- Roberto Chinicci, Oracle

1.6 Acknowledgements

During the course of the JSR we have received many excellent suggestions on the JSR mailing lists. Thanks to those people.

Chapter 2 - Caches

A cache is a place to store temporary data. It presents a Map-like API. Like a map, data is stored as values by key.

Like a map data is accessible via a key, which is a unique identifier. There can only be one entry in a cache for a given key.

The primary programming artifact is the Cache interface. The caller interacts with the Cache.

It is often assumed that data from a database is being cached. If you are using JPA then that it is certainly the case. However fundamentally anything that is expensive or time consuming to produce can be stored in a cache. Some common use cases are:

- *client side caching of Web service calls*
- *caching of expensive computations such as rendered images*
- *caching of database rows*
- *caching of NoSQL data in-process*
- *servlet response caching*
- *caching of domain object graphs*

2.1 Requirements of Keys and Values

Keys and values are not restricted in type unless specified with a generic constructor when the cache is created.

Both keys and values must be non null. An attempt to modify the cache using a null key or value will result in a NullPointerException, as will attempts to fetch from the cache using a null key.

Caches must support storeByValue. The implication of this is that there must be a way to convert keys and values to an non-object representation. *This typically means an off-heap representation capable of being sent over a network or stored in a persistent store.*

One way to achieve this is for keys and values implement Serializable however implementers may choose other methods. Any scheme to support this other than Serializable will be implementation specific and will therefore not necessarily be portable between implementations.

2.3 Generics

The Cache interface uses generics.

2.4 Looking up a Cache

A CacheManager manages Caches which are identified and accessed by name. A reference to a cache is obtained with:

```
Cache cache = cacheManager.getCache("Greg's Cache");
```

2.5 Package

The top level package name is “javax.cache”.

2.6 Lifecycle

A Cache goes through the lifecycle as defined by [javax.cache.Status](#).

A Cache can only progress in a forward direction through these statuses. Once shutdown it cannot be restarted.

2.7 Dependency on CacheManager

A cache is always obtained from a CacheManager.

Caches are bound into the lifecycle of a CacheManager. When a CacheManager shuts down all contained caches must be shut down.

Chapter 3 Cache Configuration

3.2 Resource Limits

While a fully configured cache will require a limit on the number of entries or the resources consumed by a cache, this specification does not specify cache storage topology or structure. Implementations typically place resource constraints per node and/or storage type. However because these are not defined in the specification, there is no standard way to configure them.

It is expected therefore that implementations will provide a way to place limits on resources.

Chapter 3 - Cache Operations

This chapter deals with operations on a cache, and additional APIs.

The Cache Interface is patterned on ConcurrentMap but does not extend it. This is because some methods of ConcurrentMap (and Map) have signatures that make them inefficient in a distributed environment. As an example, Map.put returns a value which in the typical case is ignored.

Cache Operations

[This](#) is the API for Cache.

Cache Entry Expiration

Entries in the cache may expire based on the last modification and/or access time as controlled by [javax.cache.CacheConfiguration.ExpiryType](#). The following table specifies what operations cause the corresponding clock to start ticking.

An invocation resets the expiry clock for the purposes of determining expiry

- * where the expiry policy is {@link CacheConfiguration.ExpiryType#ACCESSED}.
- * It has no effect on expiry based on modification.

Method	MODIFIED	ACCESSED
boolean containsKey(K key)	NO	YES
V get(K key)	NO	YES
Map<K,V> getAll(Collection<? extends K> keys)	NO	YES
V getAndPut(K key, V value)	YES	YES
V getAndRemove(K key)	NO	NO
V getAndReplace(K key, V value)	YES	YES
CacheManager getCacheManager()	NO	NO
CacheConfiguration getConfiguration()	NO	NO
String getName()	NO	NO
CacheStatistics getStatistics()	NO	NO
Iterator<Cache.Entry<K, V>> iterator()	NO	NO (but yes on it.next)
Future<V> load(K key)	YES	YES

Future<Map<K,V>> loadAll(Collection<? extends K> keys)	YES	YES
void put(K key, V value)	YES	YES
void putAll(Map<? extends K,? extends V> map)	YES	YES
boolean putIfAbsent(K key, V value)	YES (only if was absent)	YES
boolean registerCacheEntryListener(C acheEntryListener<? super K,? super V> cacheEntryListener, NotificationScope scope, boolean synchronous)	NO	NO
boolean remove(K key)	NO	NO
boolean remove(K key, V oldValue)	NO	YES (only if not matched; not removed)
void removeAll()	NO	NO
void removeAll(Collection<? extends K> keys)	NO	NO
boolean replace(K key, V value)	YES	YES
boolean replace(K key, V oldValue, V newValue)	YES (only if matched; replaced)	YES
boolean unregisterCacheEntryListen er(CacheEntryListener<?,?> cacheEntryListener)	NO	NO
<T> T unwrap(Class<T> cls)	NO	NO

Read-Through Caching

A read-through cache behaves exactly the same way as a non-read-through cache except that get() and getAll() will invoke the CacheLoader if the entries are missing.

The effect on each method invocation when a cache is in read-through mode is described in the following table:

Method	Invoke Read-Through
boolean containsKey(K key)	No

V get(K key)	Yes
Map<K,V> getAll(Collection<? extends K> keys)	Yes. Invokes loadAll
V getAndPut(K key, V value)	No
V getAndRemove(K key)	No
V getAndReplace(K key, V value)	No
Object invokeEntryProcessor(K key, EntryProcessor<K, V> entryProcessor);	No
Iterator<Cache.Entry<K, V>> iterator()	No
Future<V> load(K key)	Yes. Even when the cache is not read-through
Future<Map<K,V>> loadAll(Collection<? extends K> keys)	Yes. Uses the CacheLoader.loadAll() method. Even when the cache is not read-through
void put(K key, V value)	No
void putAll(Map<? extends K,? extends V> map)	No
boolean putIfAbsent(K key, V value)	No
boolean remove(K key)	No
boolean remove(K key, V oldValue)	No
void removeAll()	No
void removeAll(Collection<? extends K> keys)	No
boolean replace(K key, V value)	No
boolean replace(K key, V oldValue, V newValue)	No

Write-Through Caching

It is a proxy. So any mutation always invoked the CacheWriter.

We might want a removeAllNoWriteThrough to give us the capability of just zapping the cache when it is in write-through mode.

CacheLoader

The CacheLoader interface can be used to both pre-load a cache and to allow special action on a cache miss. In the later case this could be used for example to implement read-through.

TODO: need to specify how CacheLoader gets associated with a cache.

CacheEntryListener

Listeners which are sub-interfaces of CacheEntryListener may be registered which are fired when certain cache events occur:

CacheEntryCreatedListener - fired when an entry is created.

CacheEntryUpdatedListener - fired when an entry is updated.

CacheEntryRemovedListener - fired when an entry is removed.

CacheEntryReadListener - fired when an entry is read.

Listeners are registered with one of the following notification scopes:

local - events originated in this JVM

remote - events originated in a remote JVM

If a CacheLoader is registered and causes the creation or update, the event is fired.

See <https://jsr107.ci.cloudbees.com/job/jsr107api/ws/target/apidocs/javax/cache/event> for a full description.

Classloaders

See http://www.objectsource.com/j2eechapters/Ch21-ClassLoaders_and_J2EE.htm for the problem.

Chapter 4 - Cache Managers

A CacheManager is a container of and a collection of caches.

As a container it managed all aspects of the cache lifecycle.

It provides the following capabilities:

1. a means of looking up caches by name
2. acts as an XA Resource for transaction managers
3. corresponds to a caching unit for configuration purposes
4. starts all configured caches when itself is started
5. shuts down all contained caches to be shutdown when the CacheManager is shut down
6. is the integration point for clustering in distributed caches. On startup it will allocate any necessary resources and on shutdown will release them
7. provides a default cache template so that new caches can be created with meaningful defaults
8. can be either used as a singleton in which case there will be only one CacheManager or can be used with multiple CacheManagers per VM for more complex situations.
9. allows iteration, addition and deletion of caches from it.

Singleton versus Instance

Typically a given application will only require one CacheManager. For this reason CacheManagers can be used in the simplest case as Singletons.

Singleton Construction

A singleton CacheManager is constructed, and referenced, using:

```
CacheManager.getInstance()
```

In this form the configuration must be in a well-known location.

Instance Construction

Multiple instances of CacheManager may be created using:

```
new CacheManager()
```

See Chapter 7 for details of implementation bootstrapping.

Default CacheManager

For ease of use, implementers should provide default configuration so that in the absence of an implementation specific configuration file being provided by a user, a CacheManager can be created and caches added to it.

It is anticipated that Java SE will add this specification and the Reference Implementation so that this capability will be present in Java SE.

Caches must be in a CacheManager

Because a cache may need the facilities of the CacheManager such as an integration with a cache cluster, caches may not be started outside a CacheManager. They must be added first. If they are in the configuration the CacheManager will do this automatically on startup. They may also be added programmatically to a running CacheManager.

Chapter 5 - Transactions

Transactions are an optional requirement of this specification. Transactions have the meaning provided by the JTA specification^[10]. If implemented, transactions must work as specified here.

Two transaction modes are supported:

1. Global Transactions, also known as XA Transactions.
2. Local Transactions, where the transaction boundary is the CacheManager

The motivation for providing transactions is that it is often very important for caches to stay strongly consistent with databases, message queues and other XA Resources. Without transaction support cache entries will not be guaranteed to be consistent with these.

A given Cache can support one of Local or XA transactions but not both at the same time.

All or nothing

If a transactions are enabled on a cache, all operations on it must happen within a transaction context otherwise a `javax.cache.transaction.TransactionException` will be thrown.

XA Transactions

XA Transactions for caches will work as per the JTA specification and the isolation level chosen.

XA transactions require the presence of a Transaction Manager.

An attempt to operate on an XA cache outside of an XA transaction context will cause a `CacheException`.

An example using programmatic transaction control is:

```
//Get a global transaction assuming in a Java EE app server
UserTransaction utxg = jndiContext.lookup("java:comp/UserTransaction");

// start the transactions
utxg.begin();

// do work
cache1.put("key1", "value");
cache2.remove("key3");
cache3.put("key5", "value4");

// commit the transactions
utxg.commit();
```

Enlistment

Enlistment is the process by which the Transaction Manager is told that an XA Resource is going to participate

in a transaction.

The javax.transaction.xa package defines the contract between the transaction manager and the resource manager, which allows the transaction manager to enlist and delist resource objects (supplied by the resource manager driver) in JTA transactions.

Enlistments is done using TransactionManager.getTransaction().enlistResource(XAResource xaRes). The way in which a reference is obtained to a TransactionManager is not defined by the JTA specification. Java EE application servers typically use a well-known JNDI path to obtain that reference which is vendor specific.

On the first XA Resource operation start() is called by the TransactionManager.

XA is connection oriented. Caches are connection agnostic which creates an impedance mismatch.

Because we do not close connections, XAResource.end() is never called.

We expect the Transaction Managers to call end() for us before the two-phase commit protocol is started (even though this is not specified in the JTA specification). This is the behaviour of most existing TransactionManagers.

The JTA spec requires “Interleaving multiple transaction contexts using the same resource may be done by the transaction manager as long as XAResource.start and XAResource.end are invoked properly for each transaction context switch. Each time the resource is used with a different transaction, the method XAResource.end must be invoked for the previous transaction that was associated with the resource, and XAResource.start must be invoked for the current transaction context.”

Because we do not call end if the XAResource was at the CacheManager level this would imply only a single transaction could be done at a time across the CacheManager. This is a possible implementation.

It is suggested that a new XAResource is created for each transaction so that a single cache manager will have as many XAResources open as there are transactions. In this way concurrent transactions are supported. The interleaving issue is also avoided.

Another possible implementation is to create an XAResource per cache operation. This is highly concurrent but requires more calls to the expensive enlistResource() method.

Recovery

Caches must implement recovery protocols as defined by JTA. In particular XAResource.recover() must be supported.

In the case of a local in-process cache where there is no durable store, and the process has restarted thus coming up with an empty cache, it is acceptable for recover() to return an empty array of javax.transaction.xa.Xid[]. The Transaction Manager may report a Heuristic Exception in this case. *This will not prevent pending transactions being correctly recovered for other XAResources. Further because the local cache is empty any attempt to read an affected value will be a cache miss so the user logic will go elsewhere for the value.*

Local Transactions

A cache supporting transactions will support Local Transactions with the four isolation levels.

Local transactions do not require the presence of a Transaction Manager.

Local Transactions allow single phase commit across multiple cache operations against one or more caches in the same CacheManager, whether distributed or local. This lets you apply multiple changes to a CacheManager all within a single transaction. If you also want to apply changes to other resources such as a database then you need to open a transaction to them and manually handle commit and rollback to ensure consistency. (Or use XA Transactions).

For example, we have two puts to Cache A and one remove to Cache B, and 4 puts to Cache C. These can all be accommodated in a single local transaction.

The JTA API is used to control local transactions. The `javax.transaction.UserTransaction` interface provides the application the ability to control transaction boundaries programmatically.

```
//Get a transaction
UserTransaction utx = cacheManager.getUserTransaction();
// start a transaction
utx.begin();
// do work
cache1.put("key1", "value");
cache2.remove("key3");

// commit the work
utx.commit();
```

The best practice as with all local transactions is to place these steps in a try-catch block and call `rollback()` if an exception is thrown. See the JTA spec for details^[10].

An attempt to operate on a cache outside of a local transaction context will cause a `javax.cache.transaction.TransactionException`.

It is possible for a single thread to have begun a XA Transaction and a local transaction. Cache operations will then be accepted for both XA and local transaction caches because both transaction contexts exist. However the transactions are separate.

So another programmatic (bean managed in EJB language) example showing this where `cache1` and `cache2` are configured for Local Transactions and `cache3` is configured for XA Transactions would be:

```
//Get a local transaction
UserTransaction utx1 = cacheManager.getUserTransaction();
```

```
//Get a global transaction assuming in a Java EE app server
UserTransaction utxg = jndiContext.lookup("java:comp/UserTransaction");

// start the transactions
utxl.begin();
utxg.begin();

// do work
cache1.put("key1", "value");
cache2.remove("key3");
cache3.put("key5", "value4");

// commit the transactions
utxl.commit();
utxg.commit();
```

Though this works, it is not particularly useful as one transaction can succeed on commit and the other can fail.

Local Transactions has it's own exceptions that can be thrown, which are all subclasses of `CacheException`. They are:

- `TransactionException` - a general exception
- `TransactionInterruptedException` - if `Thread.interrupt()` got called while the cache was processing a transaction.
- `TransactionTimeoutException` - if a cache operation or commit is called after the transaction timeout has elapsed.

Recovery

This is relevant to XA Transactions. For local transactions if you crash before `commit()` then the changes will depend on your isolation level.

Chapter 6 - Isolation Levels

The isolation level for a cache must be set at creation time and remains immutable for the lifetime of the cache.

The isolation levels `READ_COMMITTED`, `READ_UNCOMMITTED`, `REPEATABLE_READ` and `SERIALIZABLE` are required.

READ_COMMITTED

Mutating changes are not visible to other transactions in a local cache or a distributed cache until `COMMIT` has been called.

Until then Implementations are free to either:

- *return the old copy*
- *block until commit or rollback is called*

Both approaches satisfy the `READ_COMMITTED` isolation level.

READ_UNCOMMITTED

Cache mutations are immediately visible to other transactions in a local cache or a distributed cache, subject to any propagation delay of the implementation, as if transactions were not being used.

On commit, no value changes will be observed.

On rollback, the values will be reverted to their previous values, which will of course be a visible change.

On timeout, the JTA specification states that rollback is called. So on timeout the old values will be reverted too. Exactly when the rollback occurs will be implementation dependent.

SERIALIZABLE

Mutating changes are not visible to other transactions in a local cache or a distributed cache until `COMMIT` has been called.

Further no changes to the cache made by other transactions are visible to this transaction until it completes.

The `SERIALIZABLE` isolation level offers one further protection over `REPEATABLE_READ`, protection from Phantom reads.

An alternative is to exclusively write lock a collection of keys of interest before starting your transaction. We could use `lockAll(Collection keys)`. This would create a `ReadWrite` lock. Other transactions would block until this transaction.

This behaviour could be achieved pessimistically with a `ReadWrite` lock over the entire cache or also achieved optimistically by triggering a `RollbackException` if any changes made to the keys used (for reads or writes) in

this transaction have been made.

REPEATABLE_READ

Mutating changes are not visible to other transactions in a local cache or a distributed cache until COMMIT has been called.

Further no changes to the cache made by other transactions for keys once they have or written by this transaction are visible to this transaction until it completes.

This behaviour could be achieved pessimistically with a ReadWrite lock acquired over the keys as they are used or also achieved optimistically by triggering a RollbackException if any changes made to the keys used (for writes) in this transaction have been made.

Chapter 7 - Caching Annotations

This chapter talks about the annotations defined in the Java Caching 1.0 specification. Annotations are an optional part of this specification and can be implemented as a stand-alone support library. The JavaDoc for the annotations classes should be read for detailed descriptions of functionality.

7.1 - Annotations Overview

The annotations and support classes are in the `javax.cache.annotation` package.

7.1.1 - Annotation Inheritance and Ordering

This specification defers to section 2.1 of the Common Annotations for Java specification for annotation inheritance. Order of interceptor execution with regards to annotations outside of this specification is not defined and left to the annotation support implementation.

7.1.2 - Multiple Annotations

Only one method level caching annotation can be specified on a method and only one parameter level caching annotation can be specified on a parameter. If more than one annotation is specified on a method or on a parameter then a `CacheAnnotationConfigurationException` must be thrown either at application initialization time or on method invocation.

7.1.3 - Transactions

If a cache is transactional, then a transaction context must exist when a caching annotated method is executed. If a transaction does not exist when the method is executed the Cache will throw a `CacheException`.

7.2 - Annotations

In an application, the method and configuration of processing caching annotations on classes is left to the implementation.

7.2.1 - @CacheDefaults

This is a class level annotation used to define default property values for all caching related annotations used in a class. The `cacheName`, `cacheResolverFactory`, and `cacheKeyGenerator` properties may be specified though all are optional.

If `@CacheDefaults` is specified on a class but no method level caching annotations exist then the `@CacheDefaults` annotation is ignored.

The following example shows specifying a cache named “cities” which will be used as the default cache name for the class. The `@CacheResult` annotation on the `getCity` method will use this cache name at runtime.

Code Example 7-1 @CacheDefaults Annotation Example

```
@CacheDefaults(cacheName="cities")
public class CitySource {
    @CacheResult
    public City getCity(int lat, int lon) {
        //...
    }
}
```

7.2.2 - @CacheResult

This is a method level annotation used to mark methods whose returned value is cached using a key generated from the method parameters and returned from cache on later calls with the same parameters.

The @CacheKeyParam annotation can be used to select a subset of the parameters for key generation.

Options

1. Toggle caching of null return values via the cacheNull property.
2. Optional caching and re-throwing of exceptions with their own named cache, includes the ability to only cache specific exceptions.
3. Optional skipping of the pre-execution Cache.get call, useful when the annotated method should always be executed and the returned value placed in the cache.

@CacheResult will be ignored if placed on static methods.

7.2.3 - @CachePut

This is a method level annotation used to mark methods where one of the method arguments should be stored in the cache. One parameter must be annotated with @CacheValue marking it the parameter to be cached. If no @CacheValue annotation is specified a CacheAnnotationConfigurationException must be thrown either at application initialization time or on method invocation.

The @CacheKeyParam annotation can be used to select a subset of the parameters for key generation, the @CacheValue annotated parameter is never included in key generation.

Options

1. Toggle caching of null parameter values via the cacheNull property.
2. Specify if the Cache.put call will happen before or after method execution.
3. If caching happens after invocation then an exception thrown by the annotated method can cancel the Cache.put call

@CachePut will be ignored if placed on static methods.

7.2.4 - @CacheRemoveEntry

This is a method level annotation used to mark methods where the invocation results in an entry being removed from the specified Cache.

The @CacheKeyParam annotation can be used to select a subset of the parameters for key generation.

Options

1. Specify if the Cache.remove call will happen before or after method execution
2. If removal happens after invocation then an exception thrown by the annotated method can cancel the Cache.remove call

@CacheRemoveEntry will be ignored if placed on static methods.

7.2.5 - @CacheRemoveAll

This is a method level annotation used to mark methods where the invocation results in all entries being removed from the specified Cache.

Options

1. Specify if the `Cache.removeAll` call will happen before or after method execution
2. If removal happens after invocation then an exception thrown by the annotated method can cancel the `Cache.removeAll` call

`@CacheRemoveAll` will be ignored if placed on static methods.

7.2.5 - `@CacheKeyParam`

This is a parameter level annotation used to mark parameters that are used to generate the `CachKey` via the `CacheKeyGenerator`. At execution time the values of the parameters annotated with `@CacheKeyParam` are placed in the `CacheKeyInvocationContext.getKeyParameters()` array.

Usable with `@CacheResult`, `@CachePut`, and `@CacheRemoveEntry`

7.2.6 - `@CacheValue`

This is a parameter level annotation used to mark the parameter to be cached for a method annotated with `@CachePut`. A parameter annotated with `@CachePut` will never be included in the `CacheKeyInvocationContext.getKeyParameters()` array.

Usable with `@CachePut`

7.3 - *Cache Resolution*

All of the method level annotations allow for specification of a `CacheResolverFactory` and cache name which are used to determine the Cache to interact with at runtime.

7.3.1 - Cache Name

If no cache name is specified either on the method level annotation or the `@CacheDefaults` annotation the name is generated as the following:

```
package.name.ClassName.methodName(package.ParameterType,package.ParameterType)
```

The `@CacheResult` annotation has an additional `exceptionCacheName` property, if this property is not specified there is no default exception cache name and no exception cache is used.

7.3.2 - `CacheResolverFactory`

The specified `CacheResolverFactory` must be called exactly once per annotated method to determine the `CacheResolver` to use for each execution of the annotated method. When an annotated method is executed the previously retrieved `CacheResolver` is used to determine the Cache to use based on the `CacheInvocationContext`.

If `javax.cache.annotation.CacheResolverFactory` is specified on the annotation and the `@CacheDefaults` then the default `CacheResolverFactory` logic must be used.

Default `CacheResolverFactory` Rules:

1. Get the `CacheManager` to use via `Caching.getCacheManager()`
2. Call `CacheManager.get(String)` with the cache name
3. If a Cache is not returned
 - a. Create the Cache using `CacheManager.createCacheBuilder`
4. Create a `CacheResolver` that wraps the found/created Cache and always returns the Cache.

If the `CacheResolverFactory` throws an exception the exception must be propagated up to the application code that triggered the execution of the `CacheResolverFactory`.

7.3.3 - CacheResolver

The `CacheResolver` is returned by the `CacheResolver` factory and is meant to be called on every invocation of the annotated method it was returned for, returning the `Cache` to use for that invocation.

If the `CacheResolver` throws an exception the exception must be propagated up to the application code that triggered the execution of the `CacheResolverFactory`.

7.4 - Key Generation

The `@CacheResult`, `@CachePut`, and `@CacheRemoveEntry` annotations all require a cache key to be generated and all of these annotations allow for specification of a `CacheKeyGenerator` implementation.

The specified `CacheKeyGenerator` will be called once for every annotated method invocation. Information about the annotated method and the current invocation is provided by the `CacheKeyInvocationContext`. The method parameters the developer specified to be used in the key are contained in the `CacheInvocationParameter` array returned by the `getKeyParameters()` method. A custom `CacheKeyGenerator` can use whatever information at its disposal to generate the `CacheKey`.

If `javax.cache.annotation.CacheKeyGenerator` is specified on the annotation and the `@CacheDefaults` then the default `CacheKeyGenerator` logic must be used.

Default `CacheKeyGenerator` Rules:

1. Create an `Object[]` using `CacheInvocationParameter.getValue()` from the array returned by `CacheKeyInvocationContext.getKeyParameters()`
2. Create a `CacheKey` instance that wraps the `Object[]` and uses `Arrays.deepHashCode` to calculate its `hashCode` and `Arrays.deepEquals` for comparison to other keys.

If the `CacheKeyGenerator` throws an exception the exception must be propagated up to the application code that triggered the execution of the `CacheKeyGenerator`.

7.5 - Annotation Support Classes

7.5.1 - CacheMethodDetails

Static information about a method with a caching annotation. Used by the `CacheResolverFactory` to determine the `CacheResolver` to use at runtime.

7.5.2 - CacheInvocationContext

Runtime information about the execution of a method with a caching annotation. Used by the `CacheResolver` to determine the `Cache` to use. Extends `CacheMethodDetails` so all static information is also available.

7.5.3 - CacheKeyInvocationContext

Runtime information about the execution of a method where key generation will take place (annotated with one of `@CacheResult`, `@CachePut`, or `@CacheRemoveEntry`). Used by the `CacheKeyGenerator` to create the `CacheKey` to use. Extends `CacheInvocationContext` so all standard runtime and static information is also available

7.5.4 - CacheInvocationParameter

Runtime information about a parameter for a method execution. Includes parameter annotations, position, type and value. Provided by `CacheInvocationContext` and `CacheKeyInvocationContext`

7.5.5 - CacheKey

Created by the `CacheKeyGenerator` interface the `CacheKey` is used as the key in any cache interacted with by the annotations. All `CacheKeys` must be immutable and serializable.

Chapter 8 - Container and Provider Contracts for Deployment and Bootstrapping

In Java SE environments, the `CacheManagerFactory.getCacheManager` method may require the creation of a new `CacheManager`. To do this it locates an instance of the `javax.cache.CacheManager.CacheManagerFactoryProvider` interface.

A cache provider implementation running in a Java SE environment should also act as a service provider by supplying a service provider configuration file as described in the JAR File Specification

The provider configuration file serves to export the provider implementation class to the `CacheManagerFactory` bootstrap class, positioning the provider as the factory for cache managers. The provider supplies the provider configuration file by creating a text file named `javax.cache.spi.CacheManagerFactoryProvider` and placing it in the `META-INF/services` directory of one of its JAR files. The contents of the file should be the name of the provider implementation class of the `javax.cache.CacheManager.CacheManagerFactoryProvider` interface.

Example:

A persistence vendor called ACME caching products ships a JAR called `acme.jar` that contains its cache provider implementation. The JAR includes the provider configuration file.

`acme.jar`

```
META-INF/services/javax.cache.spi.CacheManagerFactoryProvider
com/acme/cache/CacheManagerFactoryProvider.class
...
```

The contents of the `META-INF/services/javax.cache.spi.CacheManagerFactoryProvider` file is nothing more than the name of the implementation class:
`com.acme.cache.CacheManagerFactoryProvider`

Cache provider jars may be installed or made available in the same ways as other service providers, e.g. as extensions or added to the application classpath according to the guidelines in the JAR File Specification. If more than one provider jar is registered the first one found by `java.util.ServiceLoader` will be used. If no provider is found, `CacheManagerFactory.getCacheManager` will throw an `IllegalStateException`.

8.2 Use of Caching

The API provides a static means of accessing caching support through the class `javax.cache.Caching`. Examples include

```
// get the default cache manager
CacheManager defaultManager = Caching.getCacheManager();
// the default cache manager can also be fetched by name
assert defaultManager ==
Caching.getCacheManager(javax.cache.Caching.DEFAULT_CACHE_MANAGER_NAME);

//Can get a non default named CacheManager
CacheManager myCacheManager = Caching.getCacheManager("MyManager");
```

TODO: we need to talk about class loaders here particularly in the context of isolation and reclaiming resources.

Chapter 9 - Glossary

CacheManager	A container for caches, which holds references to them.
Cache	A collection of related items.
Caching Unit	A logical grouping under the control of a CacheManager
Key	A way of unambiguously identifying a unique item in a Cache.
Value	The value stored in a Cache. Any Java Object can be a value.
CacheLoader	A user-defined Class which is used to load key/value pairs into a Cache on demand.
CacheWriter	A user-defined Class which is used to write key/value pairs into a cache after a put operation.
Cache Store	A place where cache data is kept. Caches may have multiple stores.
CacheEventListener	A user-defined Class which listens to Cache events.
Eviction Policy	Method by which elements are evicted from the cache. E.g. FIFO, LFU, ...

Chapter 10 - Bibliography

- [1] Enterprise JavaBeans, v. 3.2. EJB Core Contracts and Requirements.
- [2] JSR-250: Common Annotations for the Java Platform 1.1. <http://jcp.org/en/jsr/detail?id=250>.
- [3] JSR-175: A Metadata Facility for the Java Programming language. <http://jcp.org/en/jsr/detail?id=175>.
- [4] JSR-221: JDBC 4.1 Specification. <http://java.sun.com/products/jdbc>.
- [5] Enterprise JavaBeans, Simplified API, v 3.0. <http://java.sun.com/products/ejb>.
- [6] JAR File Specification, <http://java.sun.com/j2se/1.5.0/docs/guide/jar/jar.html>.
- [7] JSR-220: Enterprise JavaBeans™, Version 3.0, Java Persistence API
- [8] JSR-336: Java SE
- [9] JSR342: Java™ Platform, Enterprise Edition 7 (Java EE 7) Specification, <http://www.jcp.org/en/jsr/detail?id=342>
- [10] JTA Specification 1.0.1, <http://www.oracle.com/technetwork/java/javaee/jta/index.html>
- [11] CDI Specification, <http://www.jcp.org/en/jsr/detail?id=299>
- [12] Unified EL is part of JSR-245, JSP2.1. See <http://www.jcp.org/en/jsr/detail?id=245>
- [13] @CacheResult and Ehcache prior art <http://code.google.com/p/ehcache-spring-annotations/>
- [14] @CacheResult with Grails prior art <http://gpc.github.com/grails-springcache/docs/manual/guide/4.%20Content%20Caching.html>
- [15] Information on properly implementing equals and hashCode (<http://java.sun.com/developer/Books/effectivejava/Chapter3.pdf>)

Appendix A - Revision History

This appendix lists the significant changes that have been made during the development of JSR107.

Early Draft 1

Created Document.

Appendix B - Open Issues

Here we list open issues

Statistics effects

In below

- Yes: indicates can modify statistics counter in RI
- **<YES?>**: indicates where RI does not, but maybe should
- **<YES?>coh** : indicates where RI does not, but maybe should (and does in Coherence implementation)

Points of reference:

Coherence does not maintain removals count

Coherence does not maintain getStartAccumulationDate()

Coherence uses double rather than float for averages

Greg's review comments in blue.

Method	Puts	Removals	Hits	Misses
boolean containsKey(K key) No			<YES?>coh	<YES?>coh
V get(K key)			Yes	Yes
Map<K,V> getAll(Collection<? extends K> keys)			Yes	Yes
V getAndPut(K key, V value) No	Yes		<YES?>	<YES?>
V getAndRemove(K key)		Yes	Yes	Yes
V getAndReplace(K key, V value)	Yes		Yes	Yes
Object invokeEntryProcessor(K key, EntryProcessor<K, V> entryProcessor); Not sure. This is really different.	<YES?>	<YES?>	<YES?>	<YES?>

Iterator<Cache.Entry<K, V>> iterator() ?		<YES?>	<YES?>	
Future<V> load(K key)				
Future<Map<K,V>> loadAll(Collection<? extends K> keys)				
void put(K key, V value)	Yes			
void putAll(Map<? extends K,? extends V> map)	Yes			
boolean putIfAbsent(K key, V value) No	Yes		<YES?>	<YES?>
boolean remove(K key)		Yes		
boolean remove(K key, V oldValue) No		Yes	<YES?>	<YES?>
void removeAll()		Yes		
void removeAll(Collection<? extends K> keys)		Yes		
boolean replace(K key, V value) N	Yes		<YES?>	<YES?>
boolean replace(K key, V oldValue, V newValue) N	Yes		<YES?>	<YES?>

Listeners

All listeners fire after the cache operation has completed and in the line of execution.

This means that a listener can not veto the operation taking place.

It can not alter the arguments to the operation.

It can not alter the return values except that it may throw a RuntimeException which will be wrapped as a

CacheEntryListenerException and propagated to the caller.

For a single operation a registered listener will be invoked at most once. If the VM(s) participating in the operation survive the operation it will fire exactly once.

Open Issue: do we stipulate that it fires in the VM where the request originated (in a distributed cache it could fire either in the calling VM or in the VM holding the primary data copy).

The table below summarises when listeners must be invoked. Conditions are on the state of the entry before the operation.

Method	CacheEntry CreatedList ener	CacheEntry ExpiredList ener	CacheEntry ReadListen er	CacheEntry RemovedLi stener	CacheEntry UpdatedLis tener
boolean containsKey(K key)					
V get(K key)			if there		
Map<K,V> getAll(Collection<? extends K> keys)			if there		
V getAndPut(K key, V value)	if missing		if there		if there
V getAndRemove(K key)			if there	if there	
V getAndReplace(K key, V value)			if there		if there
Object invokeEntryProcessor(K key, EntryProcessor<K, V> entryProcessor);	???		???	???	???
Iterator<Cache.Entry<K, V>> iterator()			if it.next().get Value()	if it.remove()	if it.next().set Value()
Future<V> load(K key)	if missing				if there
Future<Map<K,V>> loadAll(Collection<? extends K> keys)	if missing				if there
void put(K key, V value)	if missing				if there
void putAll(Map<? extends K,? extends V> map)	if missing				if there
boolean putIfAbsent(K key, V value)	if missing				
boolean remove(K key)				if there	
boolean remove(K key, V oldValue)			if there	if there && equal	

void removeAll()				if there	
void removeAll(Collection<? extends K> keys)				if there	
boolean replace(K key, V value)					if there
boolean replace(K key, V oldValue, V newValue)					if there && equal

Project Management

Resources Needed

People

Documentation Person - to format and proof read. 1 FTE for 2 weeks

TCK Team - to increase coverage and tighten up tests. 2 FTE months

Tools

?